Text generation web UI

A Gradio web UI for Large Language Models.

Its goal is to become the AUTOMATIC1111/stable-diffusion-webui of text generation.

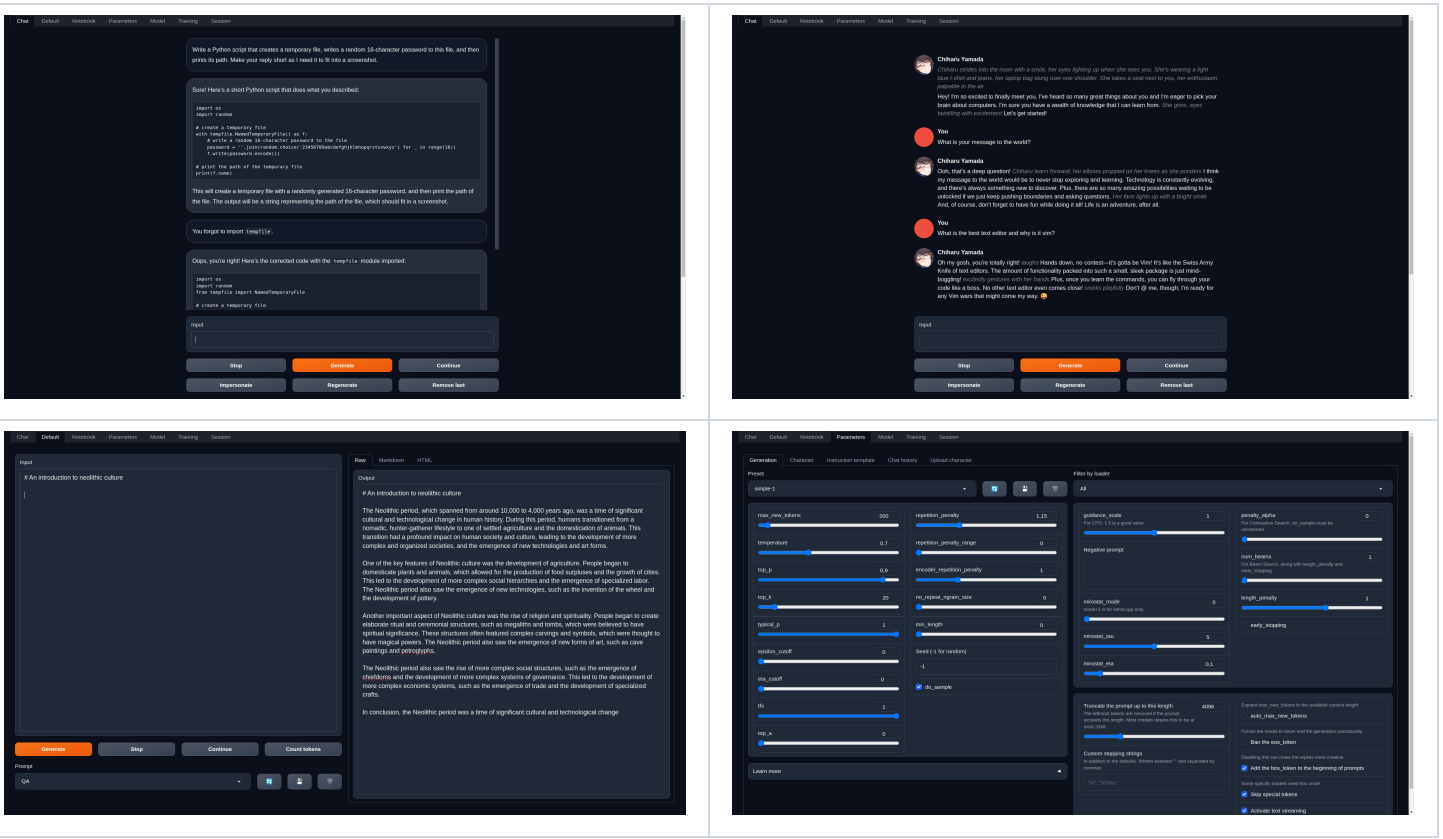

Features

- 3 interface modes: default (two columns), notebook, and chat
- Multiple model backends: transformers, llama.cpp, ExLlama, AutoGPTQ, GPTQ-for-LLaMa, ctransformers
- Dropdown menu for quickly switching between different models
- LoRA: load and unload LoRAs on the fly, train a new LoRA using QLoRA

- Precise instruction templates for chat mode, including Llama-2-chat, Alpaca, Vicuna, WizardLM, StableLM, and many others
- 4-bit, 8-bit, and CPU inference through the transformers library
- Use llama.cpp models with transformers samplers (`llamacpp_HF` loader)
- Multimodal pipelines, including LLaVA and MiniGPT-4
- Extensions framework
- Custom chat characters
- Very efficient text streaming
- Markdown output with LaTeX rendering, to use for instance with GALACTICA
- API, including endpoints for websocket streaming (see the examples)

To learn how to use the various features, check out the Documentation: https://github.com/oobabooga/text-generation-webui/tree/main/docs

Installation

One-click installers

Windows	Linux	macOS	WSL
oobabooga-windows.zip	oobabooga-linux.zip	oobabooga-macos.zip	oobabooga-wsl.zip

Just download the zip above, extract it, and double-click on "start". The web UI and all its dependencies will be installed in the same folder.

- The source codes and more information can be found here: https://github.com/oobabooga/one-click-installers
- There is no need to run the installers as admin.
- Huge thanks to @jllllll, @ClayShoaf, and @xNul for their contributions to these installers.

Manual installation using Conda

Recommended if you have some experience with the command-line.

0. Install Conda

https://docs.conda.io/en/latest/miniconda.html

On Linux or WSL, it can be automatically installed with these two commands (source):

```
curl -sL "https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh" > "Miniconda3.sh"
bash Miniconda3.sh
```

1. Create a new conda environment

```
conda create -n textgen python=3.10.9
conda activate textgen
```

2. Install Pytorch

System	GPU	Command
Linux/WSL	NVIDIA	`pip3 install torch torchvision torchaudio`
Linux/WSL	CPU only	`pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cpu`
Linux	AMD	`pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/rocm5.4.2`
MacOS + MPS	Any	`pip3 install torch torchvision torchaudio`
Windows	NVIDIA	`pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu117`
Windows	CPU only	`pip3 install torch torchvision torchaudio`

The up-to-date commands can be found here: https://pytorch.org/get-started/locally/.

3. Install the web UI

```
git clone https://github.com/oobabooga/text-generation-webui
cd text-generation-webui
pip install -r requirements.txt
```

AMD, Metal, Intel Arc, and CPUs without AVCX2

1. Replace the last command above with

```
pip install -r requirements_nocuda.txt
```

2. Manually install llama-cpp-python using the appropriate command for your hardware: Installation from PyPI.
3. AMD: Manually install AutoGPTQ: Installation.
4. AMD: Manually install ExLlama by simply cloning it into the `repositories` folder (it will be automatically compiled at runtime after that):

```
cd text-generation-webui
mkdir repositories
cd repositories
git clone https://github.com/turboderp/exllama
```

bitsandbytes on older NVIDIA GPUs

bitsandbytes >= 0.39 may not work. In that case, to use `--load-in-8bit`, you may have to downgrade like this:

- Linux: `pip install bitsandbytes==0.38.1`
- Windows: `pip install https://github.com/jllllll/bitsandbytes-windows-webui/raw/main/bitsandbytes-0.38.1-py3-none-any.whl`

Alternative: Docker

```
ln -s docker/{Dockerfile,docker-compose.yml,.dockerignore} .
cp docker/.env.example .env
# Edit .env and set TORCH_CUDA_ARCH_LIST based on your GPU model
docker compose up --build
```

- You need to have docker compose v2.17 or higher installed. See this guide for instructions.
- For additional docker files, check out this repository.

Updating the requirements

From time to time, the `requirements.txt` changes. To update, use these commands:

```
conda activate textgen
cd text-generation-webui
pip install -r requirements.txt --upgrade
```

Downloading models

Models should be placed in the `text-generation-webui/models` folder. They are usually downloaded from Hugging Face.

- Transformers or GPTQ models are made of several files and must be placed in a subfolder. Example:

```
text-generation-webui
├── models
│   ├── lmsys_vicuna-33b-v1.3
│   │   ├── config.json
│   │   ├── generation_config.json
│   │   ├── pytorch_model-00001-of-00007.bin
│   │   ├── pytorch_model-00002-of-00007.bin
│   │   ├── pytorch_model-00003-of-00007.bin
│   │   ├── pytorch_model-00004-of-00007.bin
│   │   ├── pytorch_model-00005-of-00007.bin
│   │   ├── pytorch_model-00006-of-00007.bin
│   │   ├── pytorch_model-00007-of-00007.bin
│   │   ├── pytorch_model.bin.index.json
│   │   ├── special_tokens_map.json
│   │   ├── tokenizer_config.json
│   │   └── tokenizer.model
```

- GGML/GGUF models are a single file and should be placed directly into `models`. Example:

```
text-generation-webui
├── models
│   ├── llama-13b.ggmlv3.q4_K_M.bin
```

In both cases, you can use the "Model" tab of the UI to download the model from Hugging Face automatically. It is also possible to download via the command-line with `python download-model.py organization/model` (use `--help` to see all the options).

GPT-4chan

▶ Instructions

Starting the web UI

```
conda activate textgen
cd text-generation-webui
python server.py
```

Then browse to

main text-generation-webui / README.md ↑ Top

Preview Code Blame Raw

Flag	Description
`-h` , `--help`	Show this help message and exit.
`--multi-user`	Multi-user mode. Chat histories are not saved or automatically loaded. WARNING: this is highly experimental.
`--character CHARACTER`	The name of the character to load in chat mode by default.
`--model MODEL`	Name of the model to load by default.
`--lora LORA [LORA ...]`	The list of LoRAs to load. If you want to load more than one LoRA, write the names separated by spaces.
`--model-dir MODEL_DIR`	Path to directory with all the models.
`--lora-dir LORA_DIR`	Path to directory with all the loras.
`--model-menu`	Show a model menu in the terminal when the web UI is first launched.
`--settings SETTINGS_FILE`	Load the default interface settings from this yaml file. See `settings-template.yaml` for an example. If you create a file called `settings.yaml`, this file will be loaded by default without the need to use the `--settings` flag.
`--extensions EXTENSIONS [EXTENSIONS ...]`	The list of extensions to load. If you want to load more than one extension, write the names separated by spaces.
`--verbose`	Print the prompts to the terminal.

Model loader

Flag	Description
`--loader LOADER`	Choose the model loader manually, otherwise, it will get autodetected. Valid options: transformers, autogptq, gptq-for-llama, exllama, exllama_hf, llamacpp, rwkv, ctransformers

Accelerate/transformers

Flag	Description
`--cpu`	Use the CPU to generate text. Warning: Training on CPU is extremely slow.
`--auto-devices`	Automatically split the model across the available GPU(s) and CPU.
`--gpu-memory GPU_MEMORY [GPU_MEMORY ...]`	Maximum GPU memory in GiB to be allocated per GPU. Example: `--gpu-memory 10` for a single GPU, `--gpu-memory 10 5` for two GPUs. You can also set values in MiB like `--gpu-memory 3500MiB`.
`--cpu-memory CPU_MEMORY`	Maximum CPU memory in GiB to allocate for offloaded weights. Same as above.
`--disk`	If the model is too large for your GPU(s) and CPU combined, send the remaining layers to the disk.
`--disk-cache-dir DISK_CACHE_DIR`	Directory to save the disk cache to. Defaults to `cache/`.
`--load-in-8bit`	Load the model with 8-bit precision (using bitsandbytes).
`--bf16`	Load the model with bfloat16 precision. Requires NVIDIA Ampere GPU.
`--no-cache`	Set `use_cache` to False while generating text. This reduces the VRAM usage a bit with a performance cost.
`--xformers`	Use xformer's memory efficient attention. This should increase your tokens/s.
`--sdp-attention`	Use torch 2.0's sdp attention.
`--trust-remote-code`	Set trust_remote_code=True while loading a model. Necessary for ChatGLM and Falcon.

Accelerate 4-bit

⚠️ Requires minimum compute of 7.0 on Windows at the moment.

Flag	Description
`--load-in-4bit`	Load the model with 4-bit precision (using bitsandbytes).
`--compute_dtype COMPUTE_DTYPE`	compute dtype for 4-bit. Valid options: bfloat16, float16, float32.
`--quant_type QUANT_TYPE`	quant_type for 4-bit. Valid options: nf4, fp4.
`--use_double_quant`	use_double_quant for 4-bit.

GGML/GGUF (for llama.cpp and ctransformers)

Flag	Description
`--threads`	Number of threads to use.
`--n_batch`	Maximum number of prompt tokens to batch together when calling llama_eval.

Flag	Description
`--n-gpu-layers N_GPU_LAYERS`	Number of layers to offload to the GPU. Only works if llama-cpp-python was compiled with BLAS. Set this to 1000000000 to offload all layers to the GPU.
`--n_ctx N_CTX`	Size of the prompt context.

llama.cpp

Flag	Description
`--no-mmap`	Prevent mmap from being used.
`--mlock`	Force the system to keep the model in RAM.
`--mul_mat_q`	Activate new mulmat kernels.
`--cache-capacity CACHE_CAPACITY`	Maximum cache capacity. Examples: 2000MiB, 2GiB. When provided without units, bytes will be assumed.
`--tensor_split TENSOR_SPLIT`	Split the model across multiple GPUs, comma-separated list of proportions, e.g. 18,17
`--llama_cpp_seed SEED`	Seed for llama-cpp models. Default 0 (random).
`--n_gqa N_GQA`	GGML only (not used by GGUF): Grouped-Query Attention. Must be 8 for llama-2 70b.
`--rms_norm_eps RMS_NORM_EPS`	GGML only (not used by GGUF): 5e-6 is a good value for llama-2 models.
`--cpu`	Use the CPU version of llama-cpp-python instead of the GPU-accelerated version.
`--cfg-cache`	llamacpp_HF: Create an additional cache for CFG negative prompts.

ctransformers

Flag	Description
`--model_type MODEL_TYPE`	Model type of pre-quantized model. Currently gpt2, gptj, gptneox, falcon, llama, mpt, starcoder (gptbigcode), dollyv2, and replit are supported.

AutoGPTQ

Flag	Description
`--triton`	Use triton.
`--no_inject_fused_attention`	Disable the use of fused attention, which will use less VRAM at the cost of slower inference.
`--no_inject_fused_mlp`	Triton mode only: disable the use of fused MLP, which will use less VRAM at the cost of slower inference.
`--no_use_cuda_fp16`	This can make models faster on some systems.

Flag	Description
`--desc_act`	For models that don't have a quantize_config.json, this parameter is used to define whether to set desc_act or not in BaseQuantizeConfig.
`--disable_exllama`	Disable ExLlama kernel, which can improve inference speed on some systems.

ExLlama

Flag	Description
`--gpu-split`	Comma-separated list of VRAM (in GB) to use per GPU device for model layers, e.g. `20,7,7`
`--max_seq_len MAX_SEQ_LEN`	Maximum sequence length.
`--cfg-cache`	ExLlama_HF: Create an additional cache for CFG negative prompts. Necessary to use CFG with that loader, but not necessary for CFG with base ExLlama.

GPTQ-for-LLaMa

Flag	Description
`--wbits WBITS`	Load a pre-quantized model with specified precision in bits. 2, 3, 4 and 8 are supported.
`--model_type MODEL_TYPE`	Model type of pre-quantized model. Currently LLaMA, OPT, and GPT-J are supported.
`--groupsize GROUPSIZE`	Group size.
`--pre_layer PRE_LAYER [PRE_LAYER ...]`	The number of layers to allocate to the GPU. Setting this parameter enables CPU offloading for 4-bit models. For multi-gpu, write the numbers separated by spaces, eg `--pre_layer 30 60`.
`--checkpoint CHECKPOINT`	The path to the quantized checkpoint file. If not specified, it will be automatically detected.
`--monkey-patch`	Apply the monkey patch for using LoRAs with quantized models.

DeepSpeed

Flag	Description
`--deepspeed`	Enable the use of DeepSpeed ZeRO-3 for inference via the Transformers integration.
`--nvme-offload-dir NVME_OFFLOAD_DIR`	DeepSpeed: Directory to use for ZeRO-3 NVME offloading.
`--local_rank LOCAL_RANK`	DeepSpeed: Optional argument for distributed setups.

RWKV

Flag	Description
`--rwkv-strategy RWKV_STRATEGY`	RWKV: The strategy to use while loading the model. Examples: "cpu fp32", "cuda fp16", "cuda fp16i8".
`--rwkv-cuda-on`	RWKV: Compile the CUDA kernel for better performance.

RoPE (for llama.cpp, ExLlama, and transformers)

Flag	Description
`--alpha_value ALPHA_VALUE`	Positional embeddings alpha factor for NTK RoPE scaling. Use either this or compress_pos_emb, not both.
`--rope_freq_base ROPE_FREQ_BASE`	If greater than 0, will be used instead of alpha_value. Those two are related by rope_freq_base = 10000 * alpha_value ^ (64 / 63).
`--compress_pos_emb COMPRESS_POS_EMB`	Positional embeddings compression factor. Should be set to (context length) / (model's original context length). Equal to 1/rope_freq_scale.

Gradio

Flag	Description
`--listen`	Make the web UI reachable from your local network.
`--listen-host LISTEN_HOST`	The hostname that the server will use.
`--listen-port LISTEN_PORT`	The listening port that the server will use.
`--share`	Create a public URL. This is useful for running the web UI on Google Colab or similar.
`--auto-launch`	Open the web UI in the default browser upon launch.
`--gradio-auth USER:PWD`	set gradio authentication like "username:password"; or comma-delimit multiple like "u1:p1,u2:p2,u3:p3"
`--gradio-auth-path GRADIO_AUTH_PATH`	Set the gradio authentication file path. The file should contain one or more user:password pairs in this format: "u1:p1,u2:p2,u3:p3"
`--ssl-keyfile SSL_KEYFILE`	The path to the SSL certificate key file.
`--ssl-certfile SSL_CERTFILE`	The path to the SSL certificate cert file.

API

Flag	Description
`--api`	Enable the API extension.
`--public-api`	Create a public URL for the API using Cloudfare.
`--public-api-id PUBLIC_API_ID`	Tunnel ID for named Cloudflare Tunnel. Use together with public-api option.

Flag	Description
`--api-blocking-port BLOCKING_PORT`	The listening port for the blocking API.
`--api-streaming-port STREAMING_PORT`	The listening port for the streaming API.

Multimodal

Flag	Description
`--multimodal-pipeline PIPELINE`	The multimodal pipeline to use. Examples: `llava-7b`, `llava-13b`.

Presets

Inference settings presets can be created under `presets/` as yaml files. These files are detected automatically at startup.

The presets that are included by default are the result of a contest that received 7215 votes. More details can be found here.

Contributing

If you would like to contribute to the project, check out the Contributing guidelines.

Community

- Subreddit: https://www.reddit.com/r/oobabooga/
- Discord: https://discord.gg/jwZCF2dPQN

Acknowledgment

In August 2023, Andreessen Horowitz (a16z) provided a generous grant to encourage and support my independent work on this project. I am **extremely** grateful for their trust and recognition, which will allow me to dedicate more time towards realizing the full potential of text-generation-webui.